Jaechul Roh

jrohsc.github.io · Github · Google Scholar
+1 (470) 915 - 1137 · jroh@cs.umass.edu

EDUCATION

University of Massachusetts Amherst September 2023 – May 2028 (expected)
Ph.D. in Computer Science Amherst, Massachusetts, USA
Advisor: Prof. Amir Houmansadr
GPA: 4.0/4.0

Hong Kong University of Science and Technology September 2017 – May 2023
B.Eng. in Computer Engineering, School of Engineering Clear Water Bay, Hong Kong
Final Year Thesis Advisor: Prof. Jun Zhang
*2 years of Compulsory Korean Military Duty

RESEARCH INTERESTS

My research focuses on the **Privacy & Security of AI models and agents**. Under the supervision of Prof. Amir Houmansadr, I study the trustworthiness of multimodal generative models across audio, text, and vision domains, while also collaborating closely with Prof. Eugene Bagdasarian on several projects and publications. Most recently, I completed a Summer Research Internship at Brave Software under the guidance of Dr. Ali Shahin Shamsabadi, where our work centered on advancing privacy in AI web agents.

PUBLICATIONS

Under Review

1. **SPILLage: Agentic Oversharing on the Web**
   Jaechul Roh, Eugene Bagdasarian, Hamed Haddadi, Ali Shain Shamsabadi
   Under Review

2. ***Bob's Confetti*: Phonetic Memorization Attacks in Music and Video Generation**
   Jaechul Roh, Zachary Novack, Yuefeng Peng, Niloofar Mireshghallah, Taylor Berg-Kirkpatrick, Amir Houmansadr
   Preprint at arXiv (Under Review)
   [paper] [demo page] [video]

3. **OverThink: Slowdown Attacks on Reasoning LLMs**
   Abhinav Kumar, **Jaechul Roh**, Ali Naseh, Marezna Karpinska, Mohit Iyyer, Amir Houmansadr, and Eugene Bagdasarian
   Preprint at arXiv (Under Review)
   [paper] [code]

4. **World-To-Image: Grounding Text-to-Image Generation with Agent-Driven World Knowledge**
   Moohyun Son, Jintaek Oh, Sunbin Mun, **Jaechul Roh**, Sehyun Choi
   Under Review

5. **Throttling Web Agents Using Reasoning Gates**
   Abhinav Kumar, **Jaechul Roh**, Ali Naseh, Amir Houmansadr, Eugene Bagdasarian
   Preprint at arXiv (Under Review)
   [paper] [code] [demo page]

Conference

1. **Multilingual and Multi-Accent Jailbreaking of Audio LLMs**
   Jaechul Roh, Virat Shejwalkar, Amir Houmansadr
   COLM 2025
   [paper]

2. **Backdooring Bias (B^2) into Stable Diffusion Models**
   Ali Naseh, **Jaechul Roh**, Eugene Bagdasarian, Amir Houmansadr
   USENIX Security '25
   [paper] [code]

3. **OSLO: One-Shot Label-Only Membership Inference Attacks**
   Yuefeng Peng, **Jaechul Roh**, Subhransu Maji, Amir Houmansadr
   NeurIPS 2024

[paper]

4. **Memory Triggers: Unveiling Memorization in Text-To-Image Generative Models through Word-Level Duplication**
Ali Naseh, **Jaechul Roh**, Amir Houmansadr
The 5th AAAI Workshop on Privacy-Preserving Artificial Intelligence
[paper]

5. **Robust Smart Home Face Recognition under Starving Federated Data**
Jaechul Roh, Yajun Fang
IEEE International Conference on Universal Village (IEEE UV2022)
Oral Presentation
[paper][code][slides][video]

6. **MSDT: Masked Language Model Scoring Defense in Text Domain**
Jaechul Roh, Minhao Cheng, Yajun Fang
IEEE International Conference on Universal Village (IEEE UV2022)
Oral Presentation
[paper][code][slides][video]

7. **Impact of Adversarial Training on the Robustness of Deep Neural Networks**
Jaechul Roh
2022 IEEE 5th International Conference on Information Systems and Computer Aided Education (ICISCAE)
[paper][code]

Preprints

1. **Chain-of-Code Collapse: Reasoning Failures in LLMs via Adversarial Prompting in Code Generation**
Jaechul Roh, Varun Gandhi, Shivani Anilkumar, Arin Garg
Preprint at arXiv
[paper] [code]

2. **R1dacted: Investigating Local Censorship in DeepSeek's R1 Language Model**
Ali Naseh, Harsh Chaudhari, **Jaechul Roh**, Mingshi Wu, Alina Oprea, Amir Houmansadr
Preprint at arXiv
[paper]

3. **FameBias: Embedding Manipulation Bias Attack in Text-to-Image Models**
Jaechul Roh*, Andrew Yuan*, Jinsong Mao*
Preprint at arXiv (Equal Contribution*)
[paper]

4. **Understanding (Un)Intended Memorization in Text-to-Image Generative Models**
Ali Naseh, **Jaechul Roh**, Amir Houmansadr
Preprint at arXiv
[paper]

Magazine Article

1. **Can You Trust What AI Hears (and Says)?**
Jaechul Roh
XRDS: Crossroads, The ACM Magazine for Students 31 (4), 34-39
[article]

INVITED TALKS & ACADEMIC SERVICES

Google Speech Technologies Group   **July 2025**
Paper presentation   Google DeepMind

- Presented our *"Multilingual and Multi-Accent Jailbreaking of Audio LLMs"* paper to the NFM Reading Group led by the Speech Technologies Group at Google DeepMind.
[slides]

Program Committee Member (Reviewer)

- **Main conference**: ICLR (2025, 2026)
- **Workshop**: COLM XLLM-Workshop (2025)

WORK EXPERIENCE

Brave Software   **June 2025 − September 2025**
Research Intern, Supervisor: Ali Shahin Shamsabadi   London, United Kingdom (Remote)
- Working on privacy & security of AI agents.

Super Chain AI (Conard International)   **June 2021 − August 2021**
NLP Engineer Intern, Supervisor: Cat Yung   Kowloon Bay, Hong Kong
- In charge of topic modeling and semantic analysis based on customer reviews and assigning specific semantics to the topics extracted.
- Competitors' analysis through web-scrapping customer reviews from other drop-shipping websites.

Military Service at Head Quarter of 12th Infantry Division   **July 2018 − March 2020**
Sergeant of Republic of Korea Army   Injae, Kang Won Do, Republic of Korea
- Officer Administrative Clerk Specialist
- Squad Leader of the Head Quarter

SKILLS / LANGUAGES

Programming Language: Python, C++
Languages: Korean (Native), English (Native), Chinese (Fluent)